


09. – 12.12.2019
Frankfurt am Main



Johannes Schnatterer, Cloudogu GmbH

Kubernetes-Security: 3 Dinge, die jeder Entwickler wissen sollte

#ittage
schnatterer









What about Security?





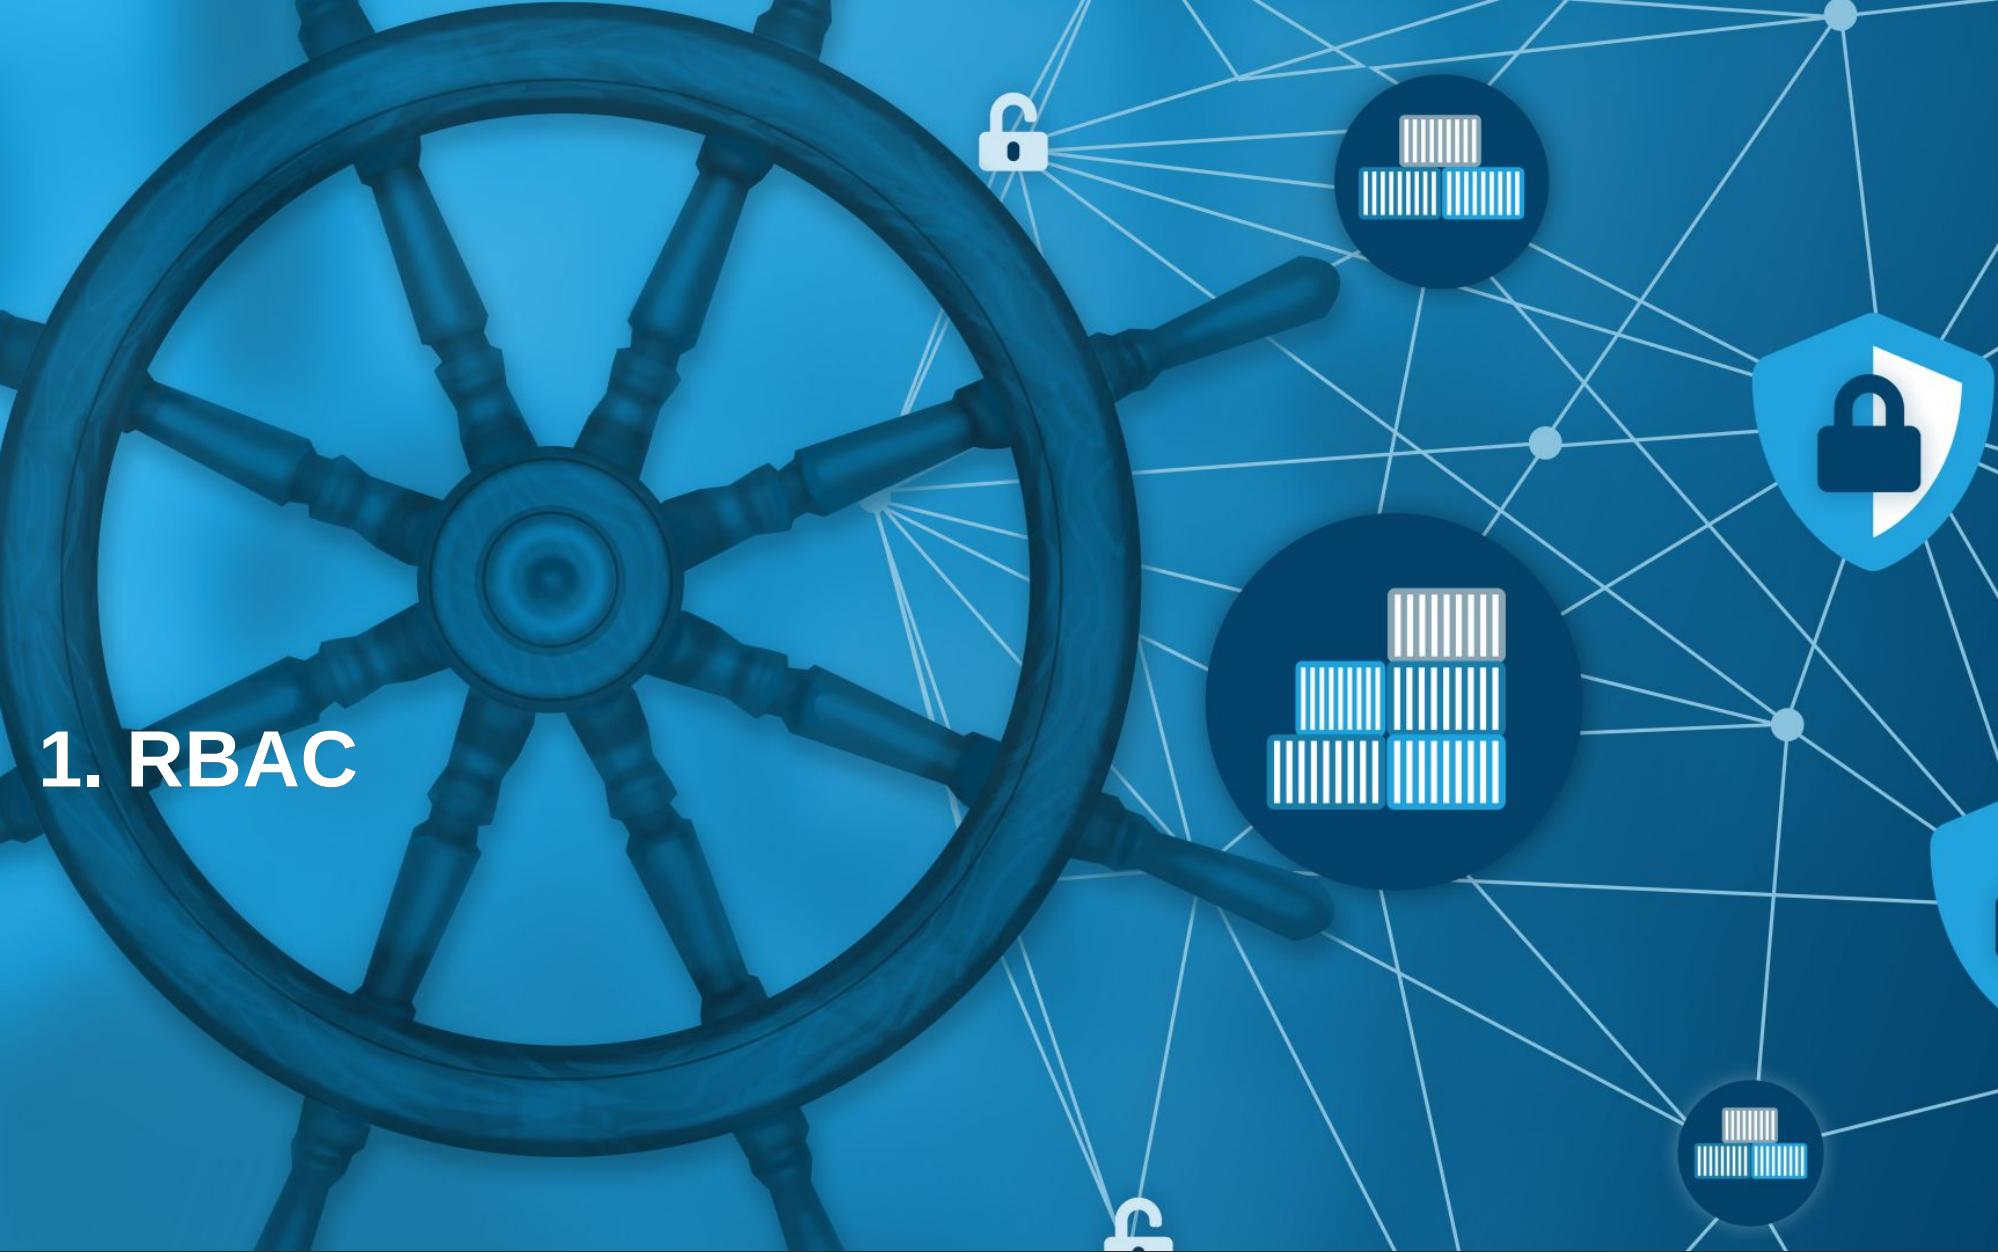
Plenty of security options

`securityContext` `runAsNonRoot` `runAsUser` `privileged` `procMount`
`allowPrivilegeEscalation` `readOnlyRootFilesystem`
`PodSecurityPolicy` `RBAC` `seccomp` `Linux Capabilities` `AppArmor`
`SELinux` `NetworkPolicy` `Falco` Open Policy Agent `gVisor` `Kata`
`Containers` `Nabla Containers` `Service Mesh` `mTLS` `KubeSec` `KubeBench`

3 things every developer should know about K8s security

- a very opinionated list of actions that make a huge difference with manageable effort
- distilled from the experience of the last years developing and operating apps on k8s

1. RBAC





🌐 <https://memegenerator.net/instance/83566913/homer-simpson-boring>

- RBAC active by default since K8s 1.6
- ... but not if you migrated!


- Every Container is mounted the token of its service account at
`/var/run/secrets/kubernetes.io/serviceaccount/token`
 - With RBAC the default service account is only authorized to read publicly accessible API info
 - ⚠️ With legacy authz the default service account is cluster admin
- You can test if your pod is authorized by executing the following in it:

```
curl --cacert /var/run/secrets/kubernetes.io/serviceaccount/ca.crt \  
-H "Authorization: Bearer $(cat /var/run/secrets/kubernetes.io/serviceaccount/token)" \  
https://${KUBERNETES_SERVICE_HOST}/api/v1/secrets
```

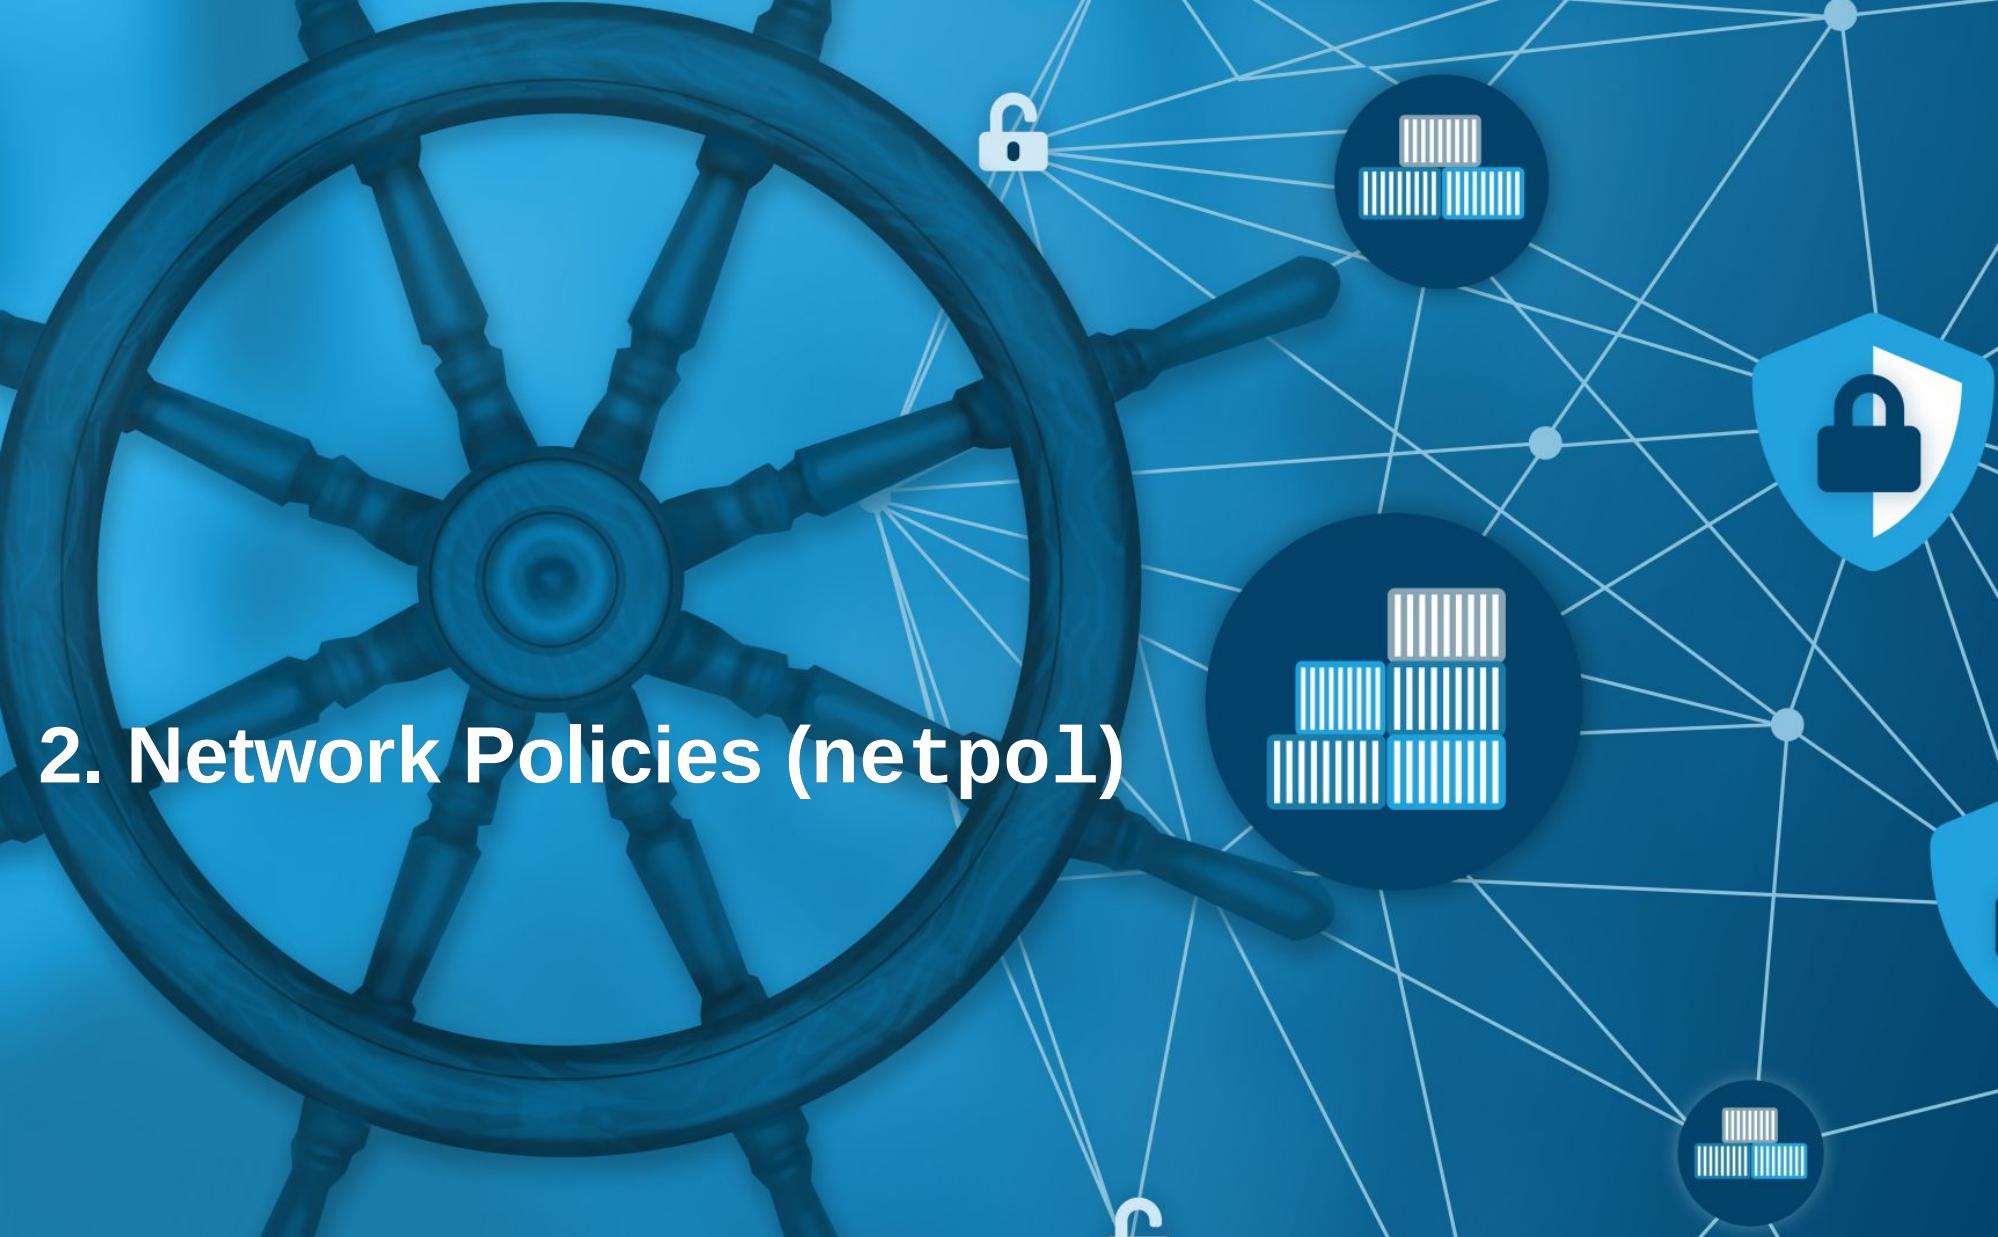
- If a pod does not need access to K8s API, mounting the token can be disabled in the pod spec:
`automountServiceAccountToken: false`



Demo



- legacy-authz
- RBAC

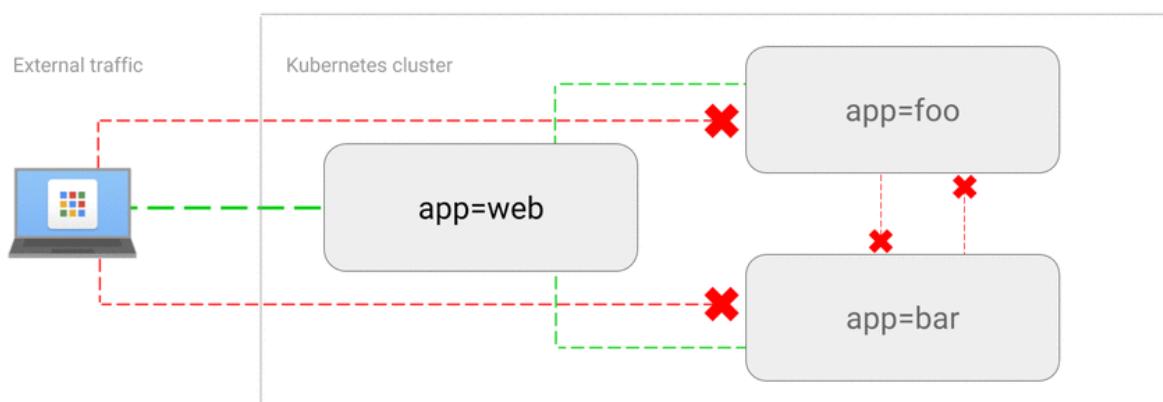


2. Network Policies (netpol)

A kind of firewall for communication between pods.

- Apply to pods (podSelector)
 - within a namespace
 - via labels
- Ingress or egress
 - to/from pods (in namespaces) or CIDRs (egress only)
 - for specific ports (optional)
- Are enforced by the CNI Plugin (e.g. Calico)
- ⚠️ No Network Policies: All traffic allowed

💡 Helpful to get started:



- 🐳 <https://github.com/ahmetb/kubernetes-network-policy-recipes>
- Securing Cluster Networking with Network Policies - Ahmet Balkan
 - <https://www.youtube.com/watch?v=3gGpMmYeEO8>
- Interactively describes what a netpol does:

```
kubectl describe netpol <name>
```

Recommendation: Whitelist ingress traffic

In every namespace except kube-system:

- Deny all ingress traffic between pods ...
- ... and then whitelist all allowed routes.

Advanced: ingress to kube-system namespace

 You might stop the apps in your cluster from working

For example, don't forget to:

- Allow external access to ingress controller
(otherwise no more external access on any cluster resource)
- Allow access to kube-dns/core-dns to every namespace
(otherwise no more service discovery by name)

Advanced: egress

- Verbose solution:
 - Deny all egress traffic between pods ...
 - ... and then whitelist all allowed routes...
 - ... repeating all ingress rules. 😕
- More pragmatic solution:
 - Allow only egress traffic within the cluster...
 - ... and then whitelist pods that need access to the internet.

Net pol pitfalls

- Don't forget to whitelist your monitoring tools (e.g. Prometheus)
- A restart of the pods might be necessary for the netpol to become effective (e.g. Prometheus)
- In order to match namespaces, labels need to be added to the namespaces, e.g.

```
kubectl label namespace/kube-system namespace=kube-system
```

- Matching both pods and namespace is only possible from k8s 1.11+
- Restricting kube-system might be more of a challenge (DNS, ingress controller)
- egress rules are more recent feature than ingress rules and seem less sophisticated
- Policies might not be supported by CNI Plugin.
Make sure to test them!
 <https://www.inovex.de/blog/test-kubernetes-network-policies/>
- On GKE: "at least 2 nodes of type n1-standard-1" are required

Limitations

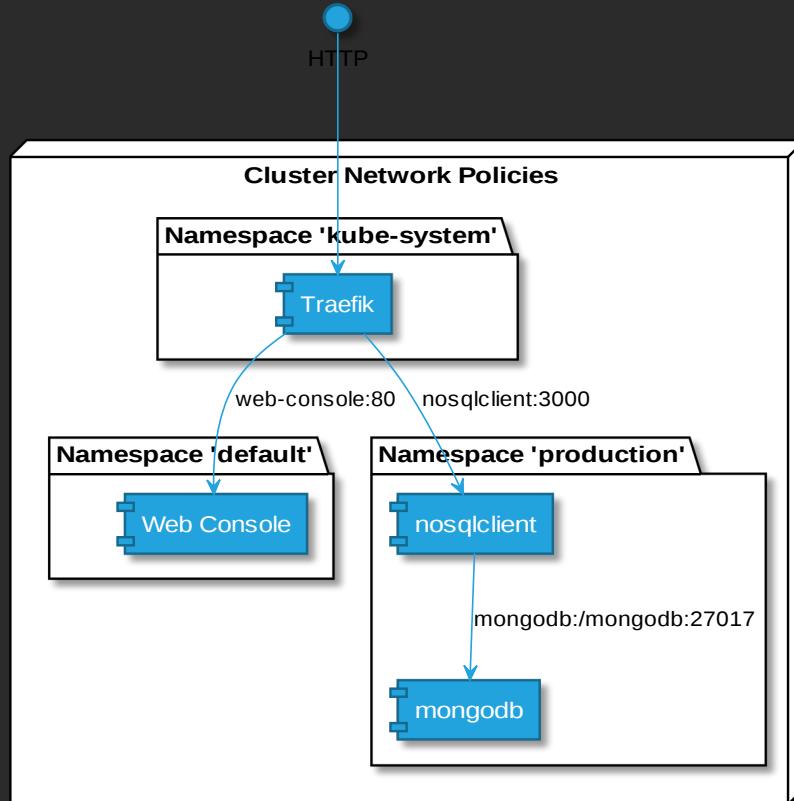
- no option for cluster-wide policies
- whitelisting egress for domain names instead of CIDRs
- filtering on L7 (e.g. HTTP or gRPC)
- netpols will not work in multi-cloud / cluster-federation scenarios

Possible solutions:

- Proprietary extensions of CNI Plugin (e.g. cilium or calico)
- Service Meshes provides similar features and work also with multiple clusters.
Service Meshes operate on L7, NetPol on L3/4
→ different strengths, support each other
 <https://istio.io/blog/2017/0.1-using-network-policy/>



Demo



- nosqlclient
- web-console



Wrap-Up: Network Policies

My recommendations:

- Definitely use DENY all ingress rule in non-kube-system namespaces
- Use with care
 - rules in kube-system
 - egress rules

3. Security Context



Defines privilege and access control settings for a Pod or Container

🌐 <https://kubernetes.io/docs/tasks/configure-pod-container/security-context/>

See also: Secure Pods - Tim Allclair

🎥 <https://www.youtube.com/watch?v=GLwmJh-j3rs>

Recommendation per Container

```
apiVersion: v1
kind: Pod
# ...
metadata:
  annotations:
    seccomp.security.alpha.kubernetes.io/pod: runtime/default
spec:
  containers:
  - name: restricted
    securityContext:
      runAsNonRoot: true
      runAsUser: 100000
      runAsGroup: 100000
      readOnlyRootFilesystem: true
      allowPrivilegeEscalation: false
      capabilities:
        drop:
        - ALL
```

There is also a securityContext on pod level, but not all of those settings cannot be applied there.

Recommendation per Container in Detail (1)

- `allowPrivilegeEscalation: false`
 - mitigates a process within the container from gaining higher privileges than its parent (the container process)
 - E.g. sudo, setuid, Kernel vulnerabilities
- `seccomp.security.alpha.kubernetes.io/pod: runtime/default`
 - Enables e.g. docker's seccomp default profile that block 44/~300 Syscalls
 - Has mitigated some Kernel vulns in the past and might in the future 🎉:
 - 🌐 <https://docs.docker.com/engine/security/non-events/>
 - no seccomp profile is also one of the findings of the k8s security audit:
 - 🌐 <https://www.cncf.io/blog/2019/08/06/open-sourcing-the-kubernetes-security-audit/>
 - `"capabilities": { "drop": ["ALL"] }`
 - Reduces attack surface
 - Drops even the default caps:
 - 🌐 <https://github.com/moby/moby/blob/master/oci/defaults.go#L14-L30>

Recommendation per Container in Detail (2)

- `runAsNonRoot: true` - Container is not started when the user is root
- `runAsUser` and `runAsGroup > 10000`
 - Reduces risk to run as user existing on host
 - In case of container escape UID/GID does not have privileges on host/filesystem
- `readOnlyRootFilesystem: true`
 - Mounts the whole file system in the container read-only. Writing only allowed in volumes.
 - Makes sure that config or code within the container cannot be manipulated.
 - It's also more efficient (no CoW).

Security context pitfalls

- `readOnlyRootFilesystem` - most applications need temp folders to write to
 - Run image locally using docker, access app (⚠️ run automated e2e/integration tests)
 - Then use `docker diff` to see a diff between container layer and image
 - and mount all folders listed there as `emptyDir` volumes in your pod
- `capabilities` - some images require capabilities
 - Start container locally with docker and `--cap-drop ALL`, then check logs for errors
 - Start again add caps as needed with e.g. `--cap-add CAP_CHOWN`, check logs for errors
 - Start again with additional caps and so forth.
 - Add all necessary caps to k8s resource
 - Alternative: Find an image of same app that does not require caps, e.g. `nginxinc/nginx-unprivileged`
- `runAsGroup` - beta from K8s 1.14. Before that defaults to GID 0 😞
 <https://github.com/kubernetes/enhancements/issues/213>

⚠️ Security context pitfalls - runAsNonRoot

- Non-root verification only supports numeric user. 😞
 - `runAsUser: 100000` in `securityContext` of pod or
 - `USER 100000` in `Dockerfile` of image.
- Some official images run as root by default.
 - Find a **trusted** image that does not run as root
 - e.g. for nginx, or postgres: 🍄 <https://hub.docker.com/r/bitnami/>
 - Derive from the original image and create your own non-root image
 - e.g. nginx: ⚙️ <https://github.com/schnatterer/nginx-unpriv>
- UID 100000 might not have permissions to read/write. Possible solutions:
 - Init Container sets permissions for PVCs
 - Wrong permissions in container → `chmod/chown` in `Dockerfile`
- Some applications require a user for UID in `/etc/passwd`
 - New image that contains a user for UID e.g. `100000` or
 - Create `/etc/passwd` with user in init container and mount into application container

Tools


Find out if your cluster adheres to these and other good security practices:

-  [controlplaneio/kubesecc](#) - manageable amount of checks
-  [Shopify/kubeaudit](#)
 - a whole lot of checks,
 - even deny all ingress and egress NetPols and AppArmor Annotations

- Be prepared for a lot of findings
- Create your own good practices



Demo





Wrap-Up: Security Context

My recommendations

- Security Context
 - Start with least privilege
 - Only differ if there's absolutely no other way
- BTW - you can enforce Security Context Settings by using Pod Security Policies.
However, those cause a lot more effort to maintain.



4 things every developer should know about K8s security?



4 things every developer should know about K8s security?

Pod Security Policies (PSP)

a cluster-level resource [...] that define a set of conditions that a pod must run with in order to be accepted into the system

🌐 <https://kubernetes.io/docs/concepts/policy/pod-security-policy/>

- can be used to enforce security context cluster-wide
 - has additional options such as block pods that try to
 - enter node's Linux namespaces (net, PID, etc.)
 - mounting the docker socket,
 - binding ports to nodes,
 - starting privileged containers
 - etc.
 - more effort than security context and different syntax as in `securityContext` 😕
- Still highly recommended!

Too much ground to cover for 45 min!



See Demo Repo on last slide

Summary

IMHO ever person working with k8s should *at least* adhere to the following:

- Enable RBAC!
- Don't allow arbitrary connection between pods.
(e.g. use Network Policies to whitelist ingresses)
- Start with least privilege for your containers:
 - Block privilege escalation via the security context of each container
 - Enable the seccomp default module via annotation of each pod
 - Try to run your containers
 - as non-root user, with UID & GID ≥ 10000 ,
 - with a read-only file system and
 - without capabilities.
- Least privilege rules can either be set per container (securityContext) or cluster-wide (PodSecurityPolicy)



Johannes Schnatterer

Cloudogu GmbH

🌐 <https://cloudogu.com/schulungen>

K8s Security series on JavaSPEKTRUM starting 05/2019

See also 🌐 <https://cloudogu.com/blog>

🐦 [@jschnatterer](#)

🐦 [@cloudogu](#)

Demo Source: 🛡️ <https://github.com/cloudogu/k8s-security-demos>